



Sandia  
National  
Laboratories

# Zero-Overhead Abstractions

## Building Flexible Vector Math Libraries with C++20

Greg von Winckel

CppCon 2025

September 15, 2025



Sandia National Laboratories is a multimission laboratory managed and operated by National Technology & Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

SAND2025-11344 C



- **The Challenge:** Writing high-performance, generic numerical code in C++ has historically involved significant compromises.
- **What we will cover:**
  - **The Problem:** "How do I make a generic vector math library able to use any appropriate data type without *depending* on that data type's library?" (Why traditional approaches fall short.)
  - **The Solution:** Modern C++ features
    - Concepts
    - Customization Point Objects
    - `tag_invoke`
  - **Alternatives:** A brief look at other approaches.
  - **A Convenient Tool:** We will provide a single-header utility library that makes writing, using, testing, and refactoring the CPOs easy.

### 3 Why This Matters



The pervasive need for real vector computing:

- **AI/ML**: The core of modern AI (deep neural networks, matrix/vector multiplications) is driving massive growth in data center power consumption.
- **Scientific Computing**: ODE/PDEs, multiphysics, numerical optimization, engineering simulation
- **Graphics & Gaming**: Fundamental to 2D/3D rendering (transformations, lighting), AR/VR.
- **Signal and Image Processing**: Computer vision, telecommunications, medical imaging.

# The Problem: A History of Compromise in C++



- **Performance vs. Abstraction:** C++ has always offered the promise of high performance, but achieving it often meant sacrificing abstraction.
  - **C-style:** Fast, but not type-safe and hard to maintain.
  - **Object-Oriented Programming:** Better abstraction, but often at the cost of performance due to virtual function calls and heap allocations.
- **CRTP:** A step in the right direction, but it's complex, verbose, and doesn't fully solve the problem of extensibility.
- **Expression Templates:** A powerful technique for optimizing mathematical expressions, but they are difficult to write and debug.
- **Other approaches:** Traits, function templates, policies, type erasure, `std::variant`

# Motivation for this Work



The Rapid Optimization Library (ROL) — <https://github.com/sandialabs/rol>

- A unique C++ library for numerical optimization with special emphasis on
  - Simulation-based (e.g. PDE-constrained) optimization
  - Stochastic optimization
  - Optimal control
  - Optimal experimental design
- Uses inheritance everywhere: vectors, objectives, constraints, algorithms, etc.
- This has some advantages, e.g. no data copy needed for user types unlike other optimizers.
- It is also problematic for supporting arbitrary elementwise functions of vectors and certain architectures such as CUDA.

## When Object Oriented is Good



- **Runtime Polymorphism**: When you need to be able to choose the type of an object at runtime, OOP *may* be the right tool for the job. It's best when you need frequent new types across module/team boundaries, especially when they can be processed by existing code without modification.
- **Clear Interfaces**: Virtual functions define clear interfaces between different parts of a program, but **DON'T** use OOP *only* because you need to define an interface.
- **Encapsulation**: OOP helps to hide implementation details and protect data from unwanted modifications.
- **SOLID Applicability**: The design needs are such that ensuring the “five guidelines” of OOP can be followed without too great of complexity.
- **Hot Take**: OOP is often unfairly maligned due to the legacy of its misuse. It is like complaining about a hammer that damaged all the nuts and screws. Use the *correct tool* for each job!

**Question:** How many have tried to implement something like `virtual` from scratch as an exercise?

# What Kinds of Algorithms Do We Want?



Take the Conjugate Gradient as a motivating example. We would like to write one function that can be used with any sequential data container that contains “real numbers.”

**Require:**  $A \in \mathbb{R}^{n \times n}$ ,  $A = A^\top$ ,  $\sigma(A) \subset (0, \infty)$ ,  $b, x_0 \in \mathbb{R}^n$

**Ensure:**  $x \approx A^{-1}b$

$$r_0 \leftarrow b - Ax_0$$

$$p_0 \leftarrow r_0$$

$$k \leftarrow 0$$

**while**  $\|r_k\| > \text{tol}$  **do**

$$\alpha_k \leftarrow \frac{r_k^\top r_k}{p_k^\top A p_k}$$

$$x_{k+1} \leftarrow x_k + \alpha_k p_k$$

$$r_{k+1} \leftarrow r_k - \alpha_k A p_k$$

$$\beta_k \leftarrow \frac{r_{k+1}^\top r_{k+1}}{r_k^\top r_k}$$

$$p_{k+1} \leftarrow r_{k+1} + \beta_k p_k$$

$$k \leftarrow k + 1$$

**end while**

Depends only on: inner\_product, scale\_in\_place, add\_in\_place, dimension, clone.

# Essential Operations for Most Real-Vector Algorithms



We can reasonably expect to need to implement five distinct kinds of operations

Operation	Mathematical Representation	Function Name
Add one vector to another	$y \leftarrow y + x, x, y \in \mathbb{R}^n$	<code>add_in_place</code>
Multiply vector by scalar	$y \leftarrow \alpha y, y \in \mathbb{R}^n, \alpha \in \mathbb{R}$	<code>scale_in_place</code>
Inner/dot product	$(x \cdot y) \in \mathbb{R}, x, y \in \mathbb{R}^n$	<code>inner_product</code>
Determine vector size	$\dim(x) = n, x \in \mathbb{R}^n$	<code>dimension</code>
Create or access scratch vector	shape-compatible temporary	<code>clone</code>

Assume `clone` performs a deep copy of shape (data may be uninitialized).

We can easily think of other functions that would be useful, but this is enough to start implementing algorithms.



## 9 OOP Implementation



Let's consider how we might an abstract base class for our (mathematical) vector

```
template<typename RealT, typename IndexT=int>
class Vector {
public:
    Vector() = default;
    virtual ~Vector() = default;
    virtual void add_in_place( const Vector& ) = 0;
    virtual void scale_in_place( RealT ) = 0;
    virtual RealT inner_product( const Vector& ) = 0;
    virtual IndexT dimension() const = 0;
    /* clone? */
};
```

- Even with this simple 5-method requirement, we are off to a troubling start.
- There is no way the clone method is flexible enough for every type of “scratch” vector we might like to have

# OOP Implementation



- In order to have runtime polymorphism, we essentially need `clone` to return a pointer to a `Vector`.
- So much for leveraging views, arenas, returning by value, etc.
- Let's *pretend* for the moment that this is not an outrageous restriction.

```
#include <memory>
```

```
template<typename RealT, typename IndexT=int>
class Vector {
public:
    Vector() = default;
    virtual ~Vector() = default;
    virtual void add_in_place( const Vector& ) = 0;
    virtual void scale_in_place( RealT ) = 0;
    virtual RealT inner_product( const Vector& ) const = 0;
    virtual IndexT dimension() const = 0;
    virtual std::unique_ptr<Vector> clone() const = 0;
};
```



Let's gloss over the matrix-vector multiplication and suppose something appropriate exists. Our function might have this signature

```
template<typename RealT, typename IndexT>
void conjugate_gradient(      Vector<RealT,IndexT>& x,
                             const Matrix<RealT,IndexT>& A,
                             const Vector<RealT,IndexT>& b,
                             RealT                               tol,
                             IndexT                               maxIter );
```



The first line of our CG algorithm already leads to frustration. We have  $r_0 \leftarrow b - Ax_0$ . We can implement this as

```
auto r = b.clone();  
auto tmp = x.clone();  
A.apply(*tmp);  
tmp->scale_in_place(-1);  
r->add_in_place(*tmp);
```

- We have to keep track of which symbols are pointers and which are not.
- CG (and many other algorithms) often include expressions that are linear combinations of vectors.
- It would be convenient to have an AXPY operation such as in  $x_{k+1} \leftarrow x_k + \alpha_k p_k$ .

## Adding an AXPY method



No problem! We'll just add another method to our base class.

```
template<typename RealT, typename IndexT=int>
class Vector {
public:
    Vector() = default;
    virtual ~Vector() = default;
    // 5 pure virtual methods
    virtual void axpy_in_place( RealT alpha, const Vector& x ) = 0;
};
```

It's just one more method derived classes must override... but it *can* be decomposed into the other 5 operations. Perhaps we shouldn't *require* it.

## Adding an optional AXPY method



Instead of making pure virtual, we can add a default implementation.

```
template<typename RealT, typename IndexT=int>
class Vector {
public:
    Vector() = default;
    virtual ~Vector() = default;
    // 5 pure virtual methods
    virtual void axpy_in_place( RealT alpha, const Vector& x ) {
        auto tmp = x.clone();
        tmp->scale_in_place(alpha);
        add_in_place(*tmp);
    }
};
```

Derived classes *should* override this for efficiency, but we do not require it.

## On the use of clone in an OOP design



- Because the `clone` method dynamically allocates memory via `std::unique_ptr`, we want to use it *sparingly*.
- We check the algorithm and see what the minimum number of “new” vectors needed are.
- If we reuse `tmp` for something else, it would be convenient to have an assignment operator.
- These distractions are getting in the way. You just want to **implement the algorithm**.
- CG is a **simple** algorithm. What if we needed arbitrary nonlinear functions?
- The **real problem** is that every time you call the algorithm, for every function call *in* the algorithm, you are repeatedly asking a question whose answer never changes: “What kind of Vector are you?”
- We used OOP to define an interface. This is misuse.
- We want to write the algorithm so that it works for **every type** that implements the operations the algorithm uses.
- There has to be **a better way!**

# Comparing Approaches



Without going into how each approach would be implemented, here's a high-level overview of the pros and cons of OOP compared with well-known alternatives

Feature / Axis	OOP (Virtual)	Traits Class	Free Function Overloads	C RTP
Performance	⚠ Runtime dispatch	✓ Zero-cost	✓ Zero-cost	✓ Zero-cost
Extensibility	✗ Inheritance only	✓ Easy (specialization)	✓ Easy via overload	✗ Tied to base class
Interoperability (STL/foreign)	✗ Needs wrappers	✓ Yes	✓ Yes	✗ No
Dynamic Dispatch	✓ Supported	✗ None	✗ None	✗ None
Discoverability	✓ Central base class	✓ Central trait	✗ ADL is scattered	✓ Localized via base
Diagnostics & Tooling	✓ Good	✗ Poor without concepts	✗ Poor unless wrapped	⚠ Error-prone
Composability	✗ Hard to mix ops	✓ Modular	✓ Modular	✓ OK in closed set
Maintenance/Scalability	✗ Inflexible	✓ Scales well	⚠ Hard to organize	⚠ Fragile base tie
User Ergonomics	✓ Simple to call	⚠ Verbose per type	⚠ Simple to write, but hard to call correctly	✓ Familiar syntax



# C++20 Concepts: A Revolution in Generic Programming



- **What are they:** Concepts are named sets of requirements on a type.
- **How they work:** The compiler checks if a type satisfies a concept at compile time.
- **Why they are useful:**
  - **Improved Error Messages:** No more pages of cryptic error messages when a type doesn't meet the requirements of a template.
  - **Easier to Read and Write:** Concepts make template code self-documenting.
  - **More Flexible:** Concepts allow us to write more generic code that can work with a wider range of types.

---

Instead of writing a base class `Vector` that has a pure virtual `add_in_place` method, we define a **concept**

```
template<typename T>
concept add_in_place_c = requires( T& y, const T& x ) {
    { add_in_place(y,x) } -> std::same_as<void>;
};
```

**Later:** The conditions for a callable `add_in_place` and type `T` to satisfy the concept.

## A better approach to clone



Ideally, we'd either like `clone` to return a smart pointer to its argument type or a type that is convertible to its argument type. Helper concept and functions:

```
template<typename T>
concept deref_c = requires( T&& x ) {
    { *std::forward<T>(x) }; // Has dereference operator
};
```

```
template<deref_c T>
auto deref_if_needed(T&& x) noexcept(noexcept(*std::forward<T>(x)))
-> decltype(*std::forward<T>(x)) {
    return *std::forward<T>(x);
}
```

```
template<typename T>
requires (!deref_c<T>)
auto deref_if_needed(T&& x) noexcept -> T&& {
    return std::forward<T>(x);
}
```

## A better approach to clone



Now we can define a concept that will ensure that we can clone an object and by passing the return value through the (possibly no-op) `deref_if_needed`, we will have *something* we can use as an argument to the over vector operation functions we need.

```
template<typename T>
concept clone_c = requires (const T& x) {
    { deref_if_needed(clone(x)) } -> std::convertible_to<T>;
};
```

We can define a far more flexible `clone` function for our type than what inheritance allows. If our vector size is known at compile time, we could stack allocate a workspace (arena) that provides access to available temporary vectors. No need to dynamically allocate vectors in algorithms if that is critical for our application.

```
auto x_cl = clone(x); auto& p = deref_if_needed(x_cl);
```

Now we don't care whether `clone` returns a pointer. No further special handling needed.



A dimension function should have an integral return type

```
template<typename T>
concept dimension_c = requires( const T& x ) {
    { dimension(x) } -> std::integral;
};

template<dimension_c T>
using dimension_t = decltype(dimension(std::declval<T>()));
```

# Inner Product Concept



For purposes of simplicity, let's define

```
template<typename T>
concept real_scalar_c = std::floating_point<T>;
```

However, we could make a real scalar concept that also works for custom types, including multiprecision and rational types.

```
template<typename T>
concept inner_product_c = requires( const T& x, const T& y ) {
    { inner_product(x,y) } -> real_scalar_c;
};
```

```
template<inner_product_c T>
using inner_product_t = decltype(inner_product(std::declval<T>(),
                                                std::declval<T>()));
```

## Scale in place Concept



This concept needs two template parameters

```
template<typename T, typename S>
concept scale_in_place_c = requires( T& y, S alpha ) {
    { scale_in_place(y,alpha) } -> std::same_as<void>;
};
```

We can later add the requirement that whatever S is, it must be convertible to `inner_product_t<T>`.

Note if the size of a scalar is larger than 3x the size of a pointer, we might prefer pass by const reference here.



- Putting it all together

```
template<typename T>
concept real_vector_c = add_in_place_c<T>    &&
                       clone_c<T>           &&
                       dimension_c<T>        &&
                       inner_product_c<T>    &&
                       scale_in_place_c<T, inner_product_t<T>>;
```

- We can now use this as the criteria to determine the set of types for which an algorithm like CG is defined.
- *However*, before we use this concept, we should *first* understand what it is actually checking for.

# Why Not Free Functions?



Free-function overloads seem simple, but at scale they cause problems:

- **Unpredictable lookup:** Overload sets depend on ADL and ordinary lookup. Small, unrelated changes (hidden friends, template args, using-decls) can change which function is called.
- **Name collisions:** Any visible `add_in_place` participates in resolution. Library “fallbacks” compete with user overloads, creating ambiguities or hijacking calls.
- **No single extension point:** Customizations are scattered across namespaces. There’s no central, discoverable place to implement or find the operation for a type.
- **Poor fallbacks & diagnostics:** Forcing ADL (e.g., with a deleted template) yields inscrutable errors when nothing matches. Valid defaults are hard to provide without interfering.
- **Composability limits:** A function name isn’t an object: you can’t pass it around as state, attach policy, or layer behavior cleanly. Boilerplate repeats per operation.

**Takeaway:** Keep the call site stable and predictable. Use a CPO (an object) plus `tag_invoke` for a single, hygienic extension point.



# Qualified vs Unqualified Names



## Qualified (explicit path)

```
std::vector<int> v;  
MyNamespace::func();  
object.member();  
ptr->member();
```

Compiler looks *exactly* where you specify

## Unqualified (no path)

```
vector<int> v;    // using std  
func();          // Where is it?  
swap(a, b);      // ???
```

Compiler must *search* for the name

When unqualified, C++ uses **TWO** search strategies...

# Two Search Strategies for Unqualified Names



## Ordinary Lookup

"The Local Detective" (Inspector Lestrade)

**When:** ALWAYS goes first

**Searches for:** ANY name  
(variables, functions, types, templates)

**Strategy:** Start local, expand outward  
current → outer → namespace → global

## ADL

"The Specialist Detective" (Sherlock Holmes)

**When:** Only for function calls  
with typed arguments

**Searches for:** ONLY functions  
(never variables or types)

**Strategy:** Check argument namespaces  
(all at once, no order)

# Two Search Strategies for Unqualified Names



## Ordinary Lookup

"The Local Detective" (Inspector Lestrade)

### The Catch:

Stops at first matching name  
(even if it's the wrong type!)

`using` declarations work

Cannot be disabled

## ADL

"The Specialist Detective" (Sherlock Holmes)

### The Power:

Finds ALL matches  
(builds overload set)

`using` declarations ignored

Disable with `(func)(args)`

# Two Search Strategies for Unqualified Names



## Ordinary Lookup

"The Local Detective" (Inspector Lestrade)

## ADL

"The Specialist Detective" (Sherlock Holmes)

### Key Insight for CPOs:

**Objects** are found by ordinary lookup (predictable)

**Functions** can be found by ADL (unpredictable)

**CPOs are objects to avoid ADL!**



## ADL would find a function...

- ✓ in same namespace as type

```
namespace math {  
    template<typename T>  
    struct Vector {  
        T x, y;  
    };  
  
    void add_in_place(Vector<T>& lhs,  
                     const Vector<T>& rhs);  
}  
  
static_assert(add_in_place_c<math::Vector<double>>>);
```



## ADL would find a function...

- ✓ in same namespace as type
- ✓ that is a friend

```
namespace physics {  
    template<typename T>  
    struct Force {  
        T magnitude;  
  
        template<typename U>  
        friend void add_in_place(      Force<U>& lhs,  
                                     const Force<U>& rhs);  
    };  
}  
  
static_assert(add_in_place_c<physics::Force<float>>());
```

# Understanding Argument Dependent Lookup (ADL)



## ADL would find a function...

- ✓ in same namespace as type
- ✓ that is a friend
- ✓ in namespace associated with template arguments

```
namespace ctrs {  
    template<typename T>  
    struct Vec {  
        std::vector<T> data;  
    };  
}  
  
namespace math {  
    template<typename U>  
    struct Point { U x, y; };  
  
    // ADL will find when trying T=Point  
    template<typename T>  
    void add_in_place(ctrs::Vec<T>& lhs,  
                     const ctrs::Vec<T>& rhs)  
}  
  
static_assert(add_in_place_c<ctrs::Vec<math::Pt>>());
```



## ADL would find a function...

- ✓ in same namespace as type
- ✓ that is a friend
- ✓ in namespace associated with template arguments
- ✗ in an unrelated namespace

```
namespace graphics {  
    template<typename T>  
    struct Color {  
        T red, green, blue;  
    };  
}  
  
namespace util { // Unrelated namespace  
    template<typename T>  
    void add_in_place( graphics::Color<T>& lhs,  
                      const graphics::Color<T>& rhs);  
}  
  
static_assert(!add_in_place_c<graphics::Color<int>>>);
```





## ADL would find a function...

- ✓ in same namespace as type
- ✓ that is a friend
- ✓ in namespace associated with template arguments
- ✗ in an unrelated namespace
- ✗ in the global namespace when type is not

```
namespace audio {  
    struct Sample {  
        float amplitude;  
    };  
}  
  
// Global scope - ADL won't find this  
void add_in_place(audio::Sample& lhs,  
                  const audio::Sample& rhs);  
  
static_assert(!add_in_place_c<audio::Sample>);
```



## ADL would find a function...

- ✓ in same namespace as type
- ✓ that is a friend
- ✓ in namespace associated with template arguments
- ✗ in an unrelated namespace
- ✗ in the global namespace when type is not
- ✗ accessible only through using declaration

```
namespace game {
    template<typename T>
    struct Player {
        T health;
    };
}

namespace ops {
    template<typename T>
    void add_in_place(game::Player<T>& lhs,
                     const game::Player<T>& rhs);
}

void test_function() {

    using ops::add_in_place; // This doesn't help ADL

    static_assert(!add_in_place_c<game::Player<int>>>);
}
```

# Understanding Argument Dependent Lookup (ADL)



Another thing that ADL will **not** find are *functors*.

```
struct add_in_place_ftor {  
    template<typename T>  
    constexpr void operator()( T& y, const T& x ) const;  
};
```

```
inline constexpr add_in_place_ftor add_in_place;
```

This fact motivated “customization point objects”



```
namespace rvf {  
struct add_in_place_ftor {  
    template<add_in_place_c T>  
    constexpr void operator() ( T& y, const T& x ) const  
    noexcept(noexcept(add_in_place(y,x))) {  
        add_in_place(y,x);  
    }  
};  
  
inline constexpr add_in_place_ftor add_in_place; // CPO  
} // namespace rvf
```



We can use slightly more compact syntax.

```
namespace rvf {  
inline constexpr struct add_in_place_ftor {  
    template<add_in_place_c T>  
    constexpr void operator() ( T& y, const T& x ) const  
    noexcept(noexcept(add_in_place(y,x))) {  
        add_in_place(y,x);  
    }  
} add_in_place; // <- ADL won't try this  
} // namespace rvf
```

# Where does the compiler search for `add_in_place`?



## ADL:

- The namespace that contains the definition of type `T`.
- Any namespace enclosing `T`'s namespace.
- If `T` is a class type, the namespaces of any base classes of `T`.
- If `T` is a template instantiation like `std::vector<int>`, the namespaces associated with the template arguments (so both `std` for `vector` and the global namespace for `int`).

## Normal unqualified lookup:

- The `rvf` namespace where the call is made.
- The global namespace.

**ISO/IEC 14882:2023** paragraph 2 of `[over.match.best]` states that if there is no unique function that is better than all other viable functions, then the call is ill-formed. The standard uses precise language here: "the call is ill-formed" means the compiler must issue a diagnostic and reject the program.

## Traditional CPO (no tag\_invoke): How it works



```
namespace rvf {  
    inline constexpr struct add_in_place_ftor {  
        template<class T>  
        constexpr auto operator()(T& y, const T& x) const  
        noexcept(noexcept(add_in_place(y, x)))  
        -> decltype(add_in_place(y, x)) {  
            // 1) Ordinary lookup ignores the object, considers functions only  
            // 2) ADL adds candidates from namespaces associated with T  
            return add_in_place(y, x); // unqualified call on purpose  
        }  
    } add_in_place; // CPO object  
}
```

- Relies on ADL to find a free function named `add_in_place` for `T`.
- You cannot safely put a fallback `add_in_place` in `rvf` because it would be found by ordinary lookup.
- Common workaround is a deleted “poison-pill” template to force ADL, which hurts diagnostics.

## Problems with traditional (non-tag\_invoke) CPOs



- **Name collisions:** Any function named like the operation visible in `rvf` participates in overload resolution via ordinary lookup. This can hijack calls or create ambiguities.
- **No safe fallback:** You can't provide a default `add_in_place` in `rvf`. Common “poison-pill” tricks (a deleted template) intentionally produce hard-to-read errors when ADL fails.
- **Unpredictable ADL sets:** Hidden friends, using-declarations, and template argument associated namespaces all influence which functions are found. Small changes elsewhere can break lookups.
- **Boilerplate per operation:** Each CPO must hand-roll layering (member vs free vs fallback), noexcept and return-type plumbing, and constraints. Easy to get subtly wrong.
- **Diagnostics:** When overload resolution fails, errors point at the wrong place (deep in the CPO), not the missing/ill-formed customization.





```
namespace rvf {  
inline constexpr struct add_in_place_ftor {  
    template<typename T>  
    constexpr auto operator() ( T& y, const T& x ) const  
    noexcept(noexcept(tag_invoke(*this,y,x))) ->  
    decltype(tag_invoke(*this,y,x)) {  
        return tag_invoke(*this,y,x);  
    }  
} add_in_place; // CPO  
} // namespace rvf
```

## What `tag_invoke` buys you



- **Namespace hygiene:** The operation name is a type (the CPO object). Only `tag_invoke` is searched by ADL, avoiding collisions with `add_in_place` in `rvf`.
- **Single extension point:** Users customize by defining `tag_invoke(tag, ...)` in their associated namespace; no need to match your operation's name.
- **Safe fallbacks:** Combine with `tag_fallback_invoke` (or a constrained fallback) to supply defaults without interfering with ADL.
- **Better constraints/diagnostics:** You constrain the CPO in terms of `tag_invoke`; failures point at missing customizations, not mysterious overload sets.
- **Same zero-cost dispatch:** Resolution still happens at compile time; No runtime overhead compared (at least when passing by reference).

## Customizing with tag\_invoke (example)



```
// In user's namespace associated with the type
namespace math {
    struct vec { /* ... */ };

    // Teach rvf::add_in_place how to handle math::vec
    inline void tag_invoke(rvf::add_in_place_ftor, vec& y, const vec& x) {
        // elementwise add y += x;
    }
}

// Then generic code just calls the CPD
template<class T>
void axpy(T& y, const T& x) {
    rvf::add_in_place(y, x); // finds tag_invoke via ADL
}
```

## References on CPOs and tag\_invoke



- Customization Point Design in C++11 and Beyond (<https://ericniebler.com/page/2/>)
- N4381 - Suggested Design for Customization Points
- P1292R0 - Customization Point Functions
- P1665R0 - Tag based customization points
- P1895R0 - tag\_invoke: A general pattern for supporting customisable functions
- Why tag\_invoke is not the solution I want | Barry's C++ Blog
- P2279R0 - We need a language mechanism for customization points
- P2547R0 - Language support for customisable functions

As of right now CPOs + tag\_invoke appears to be the *best* functional solution within the standard, but *significant* drawbacks remain.

- **Boilerplate**: Large amount of “code plumbing” needed per CPO
- **Error Prone**: Easy to make mistakes implementing. Catastrophic/inscrutable compiler errors when overload resolution fails.

No language standard solution before C++29.

## TInCuP: Tag Invoke + Customization Points



- The TInCuP library was created to mitigate the pain points of writing customization point objects with `tag_invoke`.
- Get it here: <https://github.com/sandialabs/TInCuP>
- Defines a CRTP base class that provides diagnostics and introspection for your CPO and an extensive traits class.
- Provides a Python script for generating the significant amount of boilerplate needed and another to verify existing CPOs adhere to specific format for automated refactoring as C++ evolves.
- Intended as a “bridge technology” until customizable functions are supported by the standard.



- **What you type:** The CPO is specified via JSON. A Python script generates the C++ code from jinja2 templates.

```
cpo-generator '{"cpo_name": "add_in_place", \
               "args": ["$V&: y", "const $V&: x"]}'
```

- **Generic type:** The \$ symbol indicates V is a template parameter, otherwise it would be treated as a concrete type.

**Recommendation:** Treat the generated CPO code like a Makefile generated by CMake (don't modify it).

## Using TInCuP - Simple Case



- **What you type:** The CPO is specified via JSON. A Python script generates the C++ code from jinja2 templates.

```
cpo-generator '{"cpo_name": "add_in_place", \
               "args": ["$V&: y", "const $V&: x"]}'
```

- **Generic type:** The \$ symbol indicates V is a template parameter, otherwise it would be treated as a concrete type.
- **Alternatively:** If you use Vim (VS Code and CLion integrations also available)

```
CPO add_in_place '$V&:y' 'const $V&:x'
```

**Recommendation:** Treat the generated CPO code like a Makefile generated by CMake (don't modify it).



## - What you get:

```
inline constexpr struct add_in_place_ftor final
: tincup::cpo_base<add_in_place_ftor> {
    TINCUP_CPO_TAG("add_in_place")
    inline static constexpr bool is_variadic = false;
    using tincup::cpo_base<add_in_place_ftor>::operator();
    template<typename V>
    requires tincup::invocable_c<add_in_place_ftor, V&, const V&>
    constexpr auto operator()(V& y, const V& x) const
    noexcept(tincup::nothrow_invocable_c<add_in_place_ftor, V&, const V&>)
    -> tincup::invocable_t<add_in_place_ftor, V&, const V&> {
        return tag_invoke(*this, y, x);
    }
}
```



## Using TInCuP - Simple Case



### - What you get:

```
template<typename V>
concept add_in_place_invocable_c =
    tincup::invocable_c<add_in_place_ftor, V&, const V&>;

template<typename V>
concept add_in_place_nothrow_invocable_c =
    tincup::nothrow_invocable_c<add_in_place_ftor, V&, const V&>;

template<typename V>
using add_in_place_return_t =
    tincup::invocable_t<add_in_place_ftor, V&, const V&>;

template<typename V>
using add_in_place_traits =
```



Token/Pattern	Meaning	Example input	Generated signature fragment
\$T	Generic by value	"\$T: x"	<code>template&lt;typename T&gt;(T x)</code>
\$T&	Generic lvalue reference	"\$T&: x"	<code>template&lt;typename T&gt;(T&amp; x)</code>
\$T&&	Forwarding reference	"\$T&&: x"	<code>template&lt;typename T&gt;(T&amp;&amp; x) (fwd)</code>
\$T...	Generic parameter pack (value)	"\$T...: xs"	<code>template&lt;typename... T&gt;(T... xs)</code>
\$T&...	Generic lvalue reference pack	"\$T&...: xs"	<code>template&lt;typename... T&gt;(T&amp;... xs)</code>
\$T&&...	Forwarding reference pack	"\$T&&...: xs"	<code>template&lt;typename... T&gt;(T&amp;&amp;... xs) (fwd)</code>
\$const T	Const-qualified generic	"\$const T: x"	<code>template&lt;typename T&gt;(const T x)</code>
\$const T&	Const lvalue reference	"\$const T&: x"	<code>template&lt;typename T&gt;(const T&amp; x)</code>
\$volatile T&	Volatile lvalue reference	"\$volatile T&: x"	<code>template&lt;typename T&gt;(volatile T&amp; x)</code>
Concrete	Concrete type (value/lvalue)	"int: n"	<code>int n</code>
Concrete	Concrete type (rvalue)	"std::string&&: s"	<code>std::string&amp;&amp; s</code>

# No Preprocessor Black Magic



- **Note:** The preprocessor macro is only for adding metadata and as a grep-friendly token.

```
#define TINCUP_CPO_TAG(name_str) \  
    static constexpr std::string_view name = name_str; \  
    static constexpr std::string_view qualified_name() noexcept { \  
        return "tincup::" name_str; \  
    }  
}
```

- **CRTP Base Class:** Where the “magic” is.

## Inside the CRTP Base Class cpo\_base



```
// CRTP base class for all CPOs
template<typename Derived>
struct cpo_base : public cpo_introspection<Derived>,
                  public cpo_diagnostics<Derived> {
    template<typename... Args>
    constexpr void operator()(Args&&... args) const {
        this->enhanced_fail(std::forward<Args>(args)...);
    }
}; // struct cpo_base
```

## Inside the CRTP Base Class `cpo_base`



```
// CRTP base class for all CPOs
template<typename Derived>
struct cpo_base : public cpo_introspection<Derived>,
                  public cpo_diagnostics<Derived> {
    template<typename... Args>
    constexpr void operator()(Args&&... args) const {
        this->enhanced_fail(std::forward<Args>(args)...);
    }
}; // struct cpo_base
```

**ADL note.** Because each user-defined CPO type derives from `tincup::cpo_base<Derived>`, the derived CPO's *associated classes and namespaces* include those of its base class. As a result, an unqualified call to `tag_invoke` with the CPO object as the first argument considers overloads found in namespace `tincup` during argument-dependent lookup. This enables library-defined defaults and diagnostics implemented as `tag_invoke` overloads in `tincup` to be found without extra qualification.

## Inside the CRTP Base Class `cpo_base`



```
// CRTP base class for all CPOs
template<typename Derived>
struct cpo_base : public cpo_introspection<Derived>,
                 public cpo_diagnostics<Derived> {
    template<typename... Args>
    constexpr void operator()(Args&&... args) const {
        this->enhanced_fail(std::forward<Args>(args)...);
    }
}; // struct cpo_base
```

### Consequences.

- Library-provided `tag_invoke` overloads in `tincup` participate in overload resolution automatically; user-provided overloads in the argument types' namespaces are still found via their own associated namespaces.
- The candidate set for ADL includes `tincup` by construction; keep overloads constrained to avoid unintended matches, and avoid introducing unconstrained friends in unrelated namespaces.

## Helpful Concepts and Type Aliases



```
template<typename Cp, typename...Args>
concept tag_invocable_c = requires ( const Cp& cpo, Args&&...args ) {
    { tag_invoke(cpo, std::forward<Args>(args)...) } ;
};
```

```
template<typename Cp, typename...Args>
concept invocable_c = tag_invocable_c<Cp, Args...>;
```

```
template<typename Cp, typename...Args>
concept nothrow_tag_invocable_c = requires ( const Cp& cpo, Args&&...args ) {
    { tag_invoke(cpo, std::forward<Args>(args)...) } noexcept;
};
```

```
template<typename Cp, typename...Args>
concept nothrow_invocable_c = nothrow_tag_invocable_c<Cp, Args...>;
```

```
template<typename Cp, typename...Args>
using tag_invocable_t = decltype(tag_invoke(std::declval<Cp>(),
                                             std::declval<Args>()...));
```

## Utility Class cpo\_introspection



```
template<typename Derived>
struct cpo_introspection {
    // Standard introspection using derived type
    template<typename...Args>
    static constexpr bool valid_arg_types = requires { tag_invoke(std::declval<Derived>(),
                                                                    std::declval<Args>()...);};

    template<typename...Args>
    static constexpr bool is_nothrow = requires { { tag_invoke(std::declval<Derived>(),
                                                                std::declval<Args>()...) } noexcept;};

    template<typename...Args>
    using return_type = decltype(tag_invoke(std::declval<Derived>(), std::declval<Args>()...));

    // Clean alias for return types - eliminates typename/template keywords in generated code
    template<typename...Args>
    using result_t = decltype(tag_invoke(std::declval<Derived>(), std::declval<Args>()...));

    template<template<class> typename Predicate, typename...Args>
    static constexpr bool valid_return_type = Predicate<return_type<Args...>>::value;
};
```



## Utility Class `cpo_diagnostics`



Improves compile-time error messages by checking for several common mistakes when a `tag_invoke` overload is not found. Instead of a generic error, it provides a specific hint by checking if the call would have been valid under one of the following conditions:

- **Pointer Dereferencing**: Detects if arguments are pointers or smart pointers that should have been dereferenced (e.g., using `cpo(*ptr)` instead of `cpo(ptr)`).
- **Const-Correctness**: Detects when const objects are passed to a CPO that expects a mutable (non-const) argument.
- **Argument Order**: For binary operations, it checks if the call would work by swapping the two arguments.
- **Arity**: Catches common mistakes in the number of arguments provided, such as calling a unary CPO with two arguments or vice-versa.
- **Combined Issues**: It can also detect when a combination of the above errors is present (e.g., a const pointer needs to be dereferenced and passed as non-const).
- **None Detected**: Produces a standard fallback error but still displays the full list of argument types to aid in debugging.
- **Toggle Diagnostics**: Selectively enabled or disabled with compiler definitions.



```
template <class T>
struct Vector {
    T x, y;
};

inline constexpr struct normalize_ftor final
: tincup::cpo_base<normalize_ftor> {
    TINCUP_CPO_TAG("normalize")
    using tincup::cpo_base<normalize_ftor>::operator();
    inline static constexpr bool is_variadic = false;
    template<typename V>
    requires tincup::invocable_c<normalize_ftor, const V&>
    constexpr auto operator()(const V& vec) const
    noexcept(tincup::nothrow_invocable_c<normalize_ftor, const V&>)
    -> tincup::invocable_t<normalize_ftor, const V&> {
        return tag_invoke(*this, vec);
    }
} normalize;

template<typename T>
Vector<T> tag_invoke( normalize_ftor, const Vector<T>& vec ) {
    auto norm = std::sqrt(vec.x*vec.x + vec.y*vec.y);
    return {vec.x/norm, vec.y/norm};
}

int main() {
    /* expected_error: CPO: No valid tag_invoke overload for CPO,
    but there IS a valid overload for the dereferenced arguments.
    Some arguments appear to be pointers/smart_ptrs that may need
    explicit dereferencing. Consider: cpo(*ptr) instead of cpo(ptr) */
    auto ptr = std::make_unique<Vector<double>>();
    normalize(ptr); // expect enhanced diagnostics (dereference hint)

    const Vector<double> v{4, 8};
    auto n = normalize(v); // OK
}
```

```
/app/raw.githubusercontent.com/sandialabs/TINCuP/main/single_include/tincup.hpp:760:19: error:
static assertion failed due to requirement 'always_false_v<std::unique_ptr<Vector<double>,
std::default_delete<Vector<double>>>> &>': ARGUMENT TYPES: Inspect the template instantiation
above to see actual argument types
760 |     static_assert(always_false_v<Args...>,
    |                   ^~~~~~
/app/raw.githubusercontent.com/sandialabs/TINCuP/main/single_include/tincup.hpp:909:51: note: in
instantiation of template class
'tincup::cpo_diagnostics<normalize_ftor>::show_argument_types<std::unique_ptr<Vector<double>> &>'
requested here
909 |     [[maybe_unused]] show_argument_types<Args...> display_types();
    |                   ^
/app/raw.githubusercontent.com/sandialabs/TINCuP/main/single_include/tincup.hpp:1086:11: note: in
instantiation of function template specialization
'tincup::cpo_diagnostics<normalize_ftor>::enhanced_fail<std::unique_ptr<Vector<double>> &>'
requested here
1086 |     this->enhanced_fail(std::forward<Args>(args)...);
    |     ^
<source>:54:12: note: in instantiation of function template specialization
'tincup::cpo_base<normalize_ftor>::operator()<std::unique_ptr<Vector<double>> &>' requested here
54 |     normalize(ptr); // expect enhanced diagnostics (dereference hint)
    |     ^
In file included from <source>:12:
/app/raw.githubusercontent.com/sandialabs/TINCuP/main/single_include/tincup.hpp:910:19: error:
static assertion failed due to requirement 'always_false_v<normalize_ftor>': CPO: No valid
tag_invoke overload for CPO, but there IS a valid overload for the dereferenced arguments. Some
arguments appear to be pointers/smart_ptrs that may need explicit dereferencing. Consider:
cpo(*ptr) instead of cpo(ptr)
```



By supplying string options, the TInCuP CPO generator will create both runtime and compile time call method overloads for each

```
{  
  "cpo_name": "add_in_place",  
  "args": ["$T&: y", "$const T&:x"],  
  "runtime_dispatch": {  
    "type": "string",  
    "dispatch_arg": "exec_policy",  
    "options": ["sequenced",  
                "parallel"]  
  }  
}
```



```
1 inline constexpr struct add_in_place_ftor final
2 : tincup::cpo_base<add_in_place_ftor> {
3     TINCUP_CPO_TAG("add_in_place")
4     inline static constexpr bool is_variadic = false;
5     using tincup::cpo_base<add_in_place_ftor>::operator();
6     static constexpr struct sequenced_tag {} sequenced;
7     static constexpr struct parallel_tag {} parallel;
8     static constexpr struct not_found_tag {} not_found;
9
10    inline static constexpr auto options_array
11        = tincup::string_view_array<2>{"sequenced", "parallel"};
```



```
13  template<typename T>
14  requires tincup::invocable_c<add_in_place_ftor, T&, const T&,
15  add_in_place_ftor::not_found_tag>
16  constexpr auto operator()(T& y, const T& x, std::string_view exec_policy) const
17  noexcept(tincup::nothrow_invocable_c<add_in_place_ftor, T&, const T&,
18  add_in_place_ftor::not_found_tag>) {
19      // Runtime dispatch for string
20      tincup::StringDispatch<2> dispatcher(exec_policy, options_array);
21      return dispatcher.receive([&](auto dispatch_constant) {
22          if constexpr (dispatch_constant.value < 2) {
23              if constexpr (dispatch_constant.value == 0) {
24                  return tag_invoke(*this, y, x, sequenced);
25              }
26              if constexpr (dispatch_constant.value == 1) {
27                  return tag_invoke(*this, y, x, parallel);
28              }
29          } else {
30              return tag_invoke(*this, y, x, not_found);
31          }
32      });
33  }
```



```
35 // Compile-time dispatch overloads for string
36 template<typename T>
37 requires tincup::invocable_c<add_in_place_ftor, T&, const T&, sequenced_tag>
38 constexpr auto operator()(T& y, const T& x, sequenced_tag) const
39 noexcept(tincup::nothrow_invocable_c<add_in_place_ftor, T&, const T&, sequenced_tag>) {
40     return tag_invoke(*this, y, x, sequenced);
41 }
42 template<typename T>
43 requires tincup::invocable_c<add_in_place_ftor, T&, const T&, parallel_tag>
44 constexpr auto operator()(T& y, const T& x, parallel_tag) const
45 noexcept(tincup::nothrow_invocable_c<add_in_place_ftor, T&, const T&, parallel_tag>) {
46     return tag_invoke(*this, y, x, parallel);
47 }
48 template<typename T>
49 requires tincup::invocable_c<add_in_place_ftor, T&, const T&, not_found_tag>
50 constexpr auto operator()(T& y, const T& x, not_found_tag) const
51 noexcept(tincup::nothrow_invocable_c<add_in_place_ftor, T&, const T&, not_found_tag>) {
52     return tag_invoke(*this, y, x, not_found);
53 }
54 } add_in_place;
```



<https://github.com/sandialabs/RealVectorFramework>

- **Depends on:** TlncuP to define its CPOs.
- **Core operations:** `add_in_place`, `clone`, `dimension`, `inner_product`, `scale_in_place`.
- **Advanced ops:** `axpy_in_place`, `binary_in_place`, `l2norm`, `unary_in_place`, `variadic_in_place`, `ReLU`, `softmax`.
- **Memory:** Observers and tools for creating and using memory arenas.
- **Algorithms:** CG, L-BFGS, gradient descent with bound constraints, truncated CG trust region, plus a few simple transformer models.

## Support for `std::ranges::range`



If you use an STL container that satisfies the range concept, you do not need to write *any* `tag_invoke` functions.

```
template<std::ranges::range R>
void tag_invoke( add_in_place_ftor, R& y, const R& x ) {
    std::ranges::transform(y, x, std::ranges::begin(y), std::plus<>{});
}
```

```
template<std::ranges::range R>
    requires std::copy_constructible<R>
auto tag_invoke( clone_ftor, const R& x ) {
    return R(x);
}
```

```
template<std::ranges::range R>
auto tag_invoke( dimension_ftor, const R& r ) {
    return std::ranges::size(r);
}
```



## Support for `std::ranges::range`



If you use an STL container that satisfies the range concept, you do not need to write *any* `tag_invoke` functions.

```
template<std::ranges::range R>
auto tag_invoke( inner_product_ftor, const R& x, const R& y ) {
    using value_type = std::ranges::range_value_t<R>;
    return std::inner_product(std::ranges::cbegin(x),
                              std::ranges::cend(x),
                              std::ranges::begin(y),
                              static_cast<value_type>(0));
}
```

```
template<std::ranges::range R>
void tag_invoke( scale_in_place_ftor, R& y,
                 std::ranges::range_value_t<R> alpha) {
    std::ranges::for_each(y, [alpha](auto& ye){ ye *= alpha; });
}
```



```
33 template<typename Matrix, real_vector_c Vec>
34 requires self_map_c<Matrix, Vec>
35 void conjugate_gradient( const Matrix& A,
36                          const Vec& b,
37                          Vec& x,
38                          vector_value_t<Vec> relTol = 1e-5,
39                          vector_value_t<Vec> absTol = 0,
40                          vector_size_t<Vec> maxIter = 100 ) {
41
42     auto tol = rvf::fmax(relTol * rvf::l2norm(b), absTol);
43     auto b_cl = rvf::clone(b); auto& r = rvf::deref_if_needed(b_cl);
```



```
45  A(r, x);  
46  rvf::scale_in_place(r, -1.0);  
47  rvf::add_in_place(r, b);  
48  
49  auto rho0 = inner_product(r, r);  
50  if(rvf::sqrt(rho0) < tol) return;  
51  
52  auto r_cl = rvf::clone(r); auto& p = rvf::deref_if_needed(r_cl);  
53  auto x_cl = rvf::clone(x); auto& Ap = rvf::deref_if_needed(x_cl);
```



```
55     for(vector_size_t<Vec> iter = 0; iter < maxIter; ++iter) {  
56         A(Ap, p);  
57         auto pAp = rvf::inner_product(Ap, p);  
58         auto alpha = rho0 / pAp;  
59         rvf::axpy_in_place(x, alpha, p);  
60         rvf::axpy_in_place(r, -alpha, Ap);  
61         auto rho = rvf::inner_product(r, r);  
62         if(rvf::sqrt(rho) < tol) break;  
63         auto beta = rho / rho0;  
64         rvf::scale_in_place(p, beta);  
65         rvf::add_in_place(p, r);  
66         rho0 = rho;  
67     }  
68 }
```

# Arbitrary Unary Functions $x \leftarrow f(x)$



Vim: `CPO unary_in_place '$V&:target' '$F&&:func'`

Shell: `cpo-generator '{"cpo_name": "unary_in_place", \n "args": ["$V&: x", "$F&&: func"]}'`

```
inline constexpr struct unary_in_place_ftor final
: tincup::cpo_base<unary_in_place_ftor> {
TINCUP_CPO_TAG("unary_in_place")
inline static constexpr bool is_variadic = false;
template<typename F, typename V>
requires tincup::invocable_c<unary_in_place_ftor, V&, F>
constexpr auto operator()(V& x, F&& func) const
noexcept(tincup::nothrow_invocable_c<unary_in_place_ftor, V&, F>)
-> tincup::invocable_t<unary_in_place_ftor, V&, F> {
    return tag_invoke(*this, x, std::forward<F>(func));
}
} unary_in_place;
```

# Arbitrary Unary Functions $x \leftarrow f(x)$



```
template<std::ranges::range R, typename F>
requires unary_in_place_invocable<F, std::ranges::range_value_t<R>>
void tag_invoke( unary_in_place_ftor, R& y, F&& func ) {
    std::ranges::for_each(y, [func = std::forward<F>(func)](auto& ye) mutable {
        ye = func(ye);
    });
}
```

# Arbitrary Binary Functions $x \leftarrow f(x, y)$



Vim: `CPO binary_in_place '$V&:x' '$F&&:func' 'const $V:y'`

Shell: `cpo-generator '{"cpo_name": "binary_in_place", \`  
`"args": ["$V&: x", "$F&&: func", "const $V&:y"]}]'`

```
inline constexpr struct binary_in_place_ftor final
: tincup::cpo_base<binary_in_place_ftor> {
    TINCUP_CPO_TAG("binary_in_place")
    inline static constexpr bool is_variadic = false;
    template<typename F, typename V>
    requires tincup::invocable_c<binary_in_place_ftor, V&, F, const V>
    constexpr auto operator()(V& x, F&& func, const V y) const
        noexcept(tincup::nothrow_invocable_c<binary_in_place_ftor, V&, F, const V>)
        -> tincup::invocable_t<binary_in_place_ftor, V&, F, const V> {
        return tag_invoke(*this, x, std::forward<F>(func), y);
    }
} binary_in_place;
```

## Arbitrary Binary Functions $x \leftarrow f(x, y)$



```
template<std::ranges::range R, typename F>
requires binary_in_place_invocable<F, std::ranges::range_value_t<R>>
void tag_invoke( binary_in_place_ftor, R& x, F&& func, const R& y ) {
    std::ranges::transform(y, x, std::ranges::begin(y), std::forward<F>(func));
}
```



# Arbitrary Variadic Functions $x \leftarrow f(x, y, z, \dots)$



Vim: `CPO variadic_in_place '$V&:x' '$F&&:func' 'const $Vs&...:args'`

Shell: `cpo-generator '{"cpo_name": "variadic_in_place", \`  
`"args": ["$V&: x", "$F&&: func", "const $Args&...:args"]}'`

```
inline constexpr struct variadic_in_place_ftor final
: tincup::cpo_base<variadic_in_place_ftor> {
    TINCUP_CPO_TAG("variadic_in_place")
    inline static constexpr bool is_variadic = true;
    template<typename F, typename V, typename... Vs>
    requires tincup::invocable_c<variadic_in_place_ftor, V&, F, const Vs&...>
    constexpr auto operator()(V& x, F&& func, const Vs&... args) const
    noexcept(tincup::nothrow_invocable_c<variadic_in_place_ftor, V&, F, const Vs&...>)
        -> tincup::invocable_t<variadic_in_place_ftor, V&, F, const Vs&...> {
        return tag_invoke(*this, x, std::forward<F>(func), args...);
    }
} variadic_in_place;
```

## Arbitrary Variadic Functions $x \leftarrow f(x, y, z, \dots)$



```
template<typename F, typename T, typename...Args>
concept variadic_in_place_invocable =
    std::convertible_to<std::invoke_result_t<F,T,Args...>,T> &&
    (std::is_same_v<T,Args> && ...);

template<class OutputIt, class F, class FirstInput, class...RestInput>
OutputIt variadic_transform( OutputIt    out_begin,
                           OutputIt    out_end,
                           F&&          func,
                           FirstInput   first,
                           RestInput... rest ) {
    while(out_begin != out_end) {
        *out_begin++ = std::forward<F>(func)(it_inc(first), it_inc(rest)...);
    }
    return out_begin;
}
```

## Arbitrary Variadic Functions $x \leftarrow f(x, y, z, \dots)$



```
template<std::ranges::range R, typename F, typename...Args>
requires variadic_in_place_invocable<F, std::ranges::range_value_t<R>,
                                     std::ranges::range_value_t<Args>...>
void tag_invoke( variadic_in_place_ftor, R& x, F&& func, const Args&... args ) {
    return variadic_transform(
        std::ranges::begin(x),
        std::ranges::end(x),
        std::forward<F>(func),
        std::ranges::begin(args)...
    );
}
```

## Gradient Descent with Bound Constraints



*// Objective function concept*

```
template<typename F, typename Vec>
```

```
concept objective_function_c = requires(const F& f, const Vec& x, Vec& grad) {  
    { f.value(x) } -> std::convertible_to<vector_value_t<Vec>>;  
    { f.gradient(grad, x) } -> std::same_as<void>;  
};
```

*// Bound constraints representation*

```
template<real_vector_c Vec>
```

```
struct bound_constraints {
```

```
    Vec lower, upper;
```

*// Project x onto [lower, upper] bounds*

```
void project(Vec& x) const {
```

```
    binary_in_place(x, [](auto xi, auto li) { return rvf::fmax(xi, li); }, lower)
```

```
    binary_in_place(x, [](auto xi, auto ui) { return rvf::fmin(xi, ui); }, upper)
```

```
}
```

```
};
```



## - RealVectorFramework:

- Implement execution policies
- Add more algorithms
- Recreate ROL with CPOs instead of inheritance

## - TInCuP:

- Review P2547R0 and P2279R0
- Great ideas in these papers, but no implementation given yet
- It should not be too difficult to write a clang extension that uses an approach like TInCuP under the hood
- Distribute it and seek independent testing.
- A compelling possibility for C++29
- In the meantime, use TInCuP as a bridge technology with the promise of easy refactoring should a new language mechanism be introduced.

`http://github.com/sandialabs/TInCuP`

